

# Lab Report 3: Binary Classification with Neural Networks

## Circles Dataset, PyTorch Artificial Neural Networks

Name	Chetana Thapa
ROLL NO	ACE079BCT026
Lab No.	3
Subject	Deep Learning / Machine Learning
Topic	Binary Classification with Artificial Neural Networks
Framework	PyTorch
Dataset	<code>circles_binary_classification.csv</code> (Circles Dataset)

## Objective

Using the circles dataset, **build, train, evaluate, and compare multiple PyTorch Artificial Neural Networks (ANNs)** for binary classification. We will systematically increase model complexity and compare performance, culminating in a discussion of findings.

## Theoretical Background

### 1. Binary Classification

Binary classification is a supervised learning task where each input must be assigned to one of **two classes** (labeled 0 or 1). The goal is to learn a decision boundary that separates the two classes in feature space.

### 2. Artificial Neural Networks (ANNs)

An ANN is a computational model loosely inspired by biological neural networks. It consists of:

- **Input Layer:** Receives raw feature values (here:  $X_1, X_2$ )
- **Hidden Layers:** Learn intermediate representations through weighted connections and activation functions
- **Output Layer:** Produces a single logit (raw score) for binary classification

Each neuron computes:  $z = w^T x + b$ , and optionally applies an activation:  $a = \sigma(z)$

### 3. Activation Functions

- **Linear (None):**  $f(x) = x$  — No non-linearity; unable to learn non-linear decision boundaries.

- **ReLU:**  $f(x) = \max(0, x)$  — Introduces non-linearity; allows learning of complex patterns like circles.
- **Sigmoid:**  $\sigma(x) = \frac{1}{1+e^{-x}}$  — Squashes output to  $(0, 1)$ ; used for final probability.

## 4. Loss Function — BCEWithLogitsLoss

$$L = -\frac{1}{N} \sum_{i=1}^N [y_i \log(\hat{p}_i) + (1 - y_i) \log(1 - \hat{p}_i)]$$

where  $\hat{p}_i = \sigma(z_i)$ . PyTorch's `nn.BCEWithLogitsLoss` combines sigmoid + BCE in one numerically stable operation.

## 5. Optimization — SGD

Stochastic Gradient Descent updates parameters by:

$$\theta \leftarrow \theta - \alpha \nabla_{\theta} L$$

where  $\alpha$  is the learning rate. We use full-batch gradient descent here.

## 6. Decision Boundary

The set of points where the model outputs  $\hat{p}=0.5$  (i.e., logit = 0). Visualizing this reveals how well the model has learned the structure of the data.

## Section 1: Imports & Setup

```
# — Standard library & data tools
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import matplotlib.colors as mcolors
from sklearn.datasets import make_circles
from sklearn.model_selection import train_test_split

# — PyTorch
import torch
import torch.nn as nn
from torch import optim

# — Reproducibility
torch.manual_seed(42)
np.random.seed(42)

print(f"PyTorch version : {torch.__version__}")
print(f"NumPy version   : {np.__version__}")
print(f"Pandas version  : {pd.__version__}")
```

```
PyTorch version : 2.10.0+cpu
NumPy version   : 2.4.0
Pandas version  : 2.3.3
```

## Section 2: Data Retrieval & Inspection

### Theory

The **circles dataset** is a synthetic 2D dataset where one class forms an inner ring and the other forms an outer ring. It is a classic benchmark for testing whether a model can learn **non-linear decision boundaries**—linear classifiers completely fail here.

We attempt to load from CSV first; if unavailable, we generate it programmatically using `sklearn.datasets.make_circles` (identical parameters).

```
# Load or Generate Dataset
import os

CSV_PATH = 'circles_binary_classification.csv'

if os.path.exists(CSV_PATH):
    df = pd.read_csv(CSV_PATH)
    print(f"□ Loaded from CSV: {CSV_PATH}")
else:
    print("△ CSV not found – generating dataset with
sklearn.make_circles (same parameters)")
    X_raw, y_raw = make_circles(n_samples=1000, noise=0.03,
random_state=42)
    df = pd.DataFrame(X_raw, columns=['X1', 'X2'])
    df['label'] = y_raw
    df.to_csv(CSV_PATH, index=False)           # save for
reproducibility
    print(f"□ Dataset saved to {CSV_PATH}")

print(f"\nDataset shape: {df.shape}")
print(f"Columns      : {list(df.columns)}")

□ Loaded from CSV: circles_binary_classification.csv

Dataset shape: (1000, 3)
Columns      : ['X1', 'X2', 'label']

# First Five Rows
print("=" * 45)
print("          HEAD  (first 5 rows)")
print("=" * 45)
df.head()
```

```

=====
      HEAD  (first 5 rows)
=====

      X1      X2  label
0  0.754246  0.231481    1
1 -0.756159  0.153259    1
2 -0.815392  0.173282    1
3 -0.393731  0.692883    1
4  0.442208 -0.896723    0

# Statistical Summary
print("=" * 45)
print("      DESCRIBE  (statistical summary)")
print("=" * 45)
df.describe()

=====
      DESCRIBE  (statistical summary)
=====

      X1          X2      label
count  1000.000000  1000.000000  1000.00000
mean   -0.000448   -0.000804   0.50000
std    0.639837   0.641156   0.50025
min   -1.059502   -1.067768   0.00000
25%   -0.619251   -0.612176   0.00000
50%   0.008762   -0.003949   0.50000
75%   0.621933   0.624822   1.00000
max   1.033712   1.036004   1.00000

# Class Distribution
print("Class distribution (label counts):")
print(df['label'].value_counts())
print(f"\nClass balance:\n{df['label'].value_counts(normalize=True).round(3).to_dict()}")
print("\nMissing values per column:")
print(df.isnull().sum())

Class distribution (label counts):
label
1    500
0    500
Name: count, dtype: int64

Class balance: {1: 0.5, 0: 0.5}

Missing values per column:
X1      0
X2      0

```

```
label    0
dtype: int64
```

## Section 3: Data Cleaning & Feature Design

### Theory

**Data cleaning** ensures the dataset is free from:

- Missing / NaN values (would break tensor operations)
- Duplicate records (would bias training)
- Incorrect dtypes

**Feature engineering** here is minimal — X1 and X2 are already the correct 2D Cartesian coordinates. The target `label` (0 or 1) is converted to `float32` because `BCEWithLogitsLoss` requires float targets.

**Dtype convention in PyTorch:**

- Features → `torch.float32` (default floating point)
- Labels → `torch.float32` (required by `BCEWithLogitsLoss`)

```
# Data Cleaning
print("--- Before cleaning ---")
print(f"Shape: {df.shape}")
print(f"NaN values:\n{df.isnull().sum()}")
print(f"Duplicate rows: {df.duplicated().sum()}")

# Drop duplicates & nulls (if any)
df = df.drop_duplicates().dropna().reset_index(drop=True)

print("\n--- After cleaning ---")
print(f"Shape: {df.shape}")
print("Dataset is clean.")

--- Before cleaning ---
Shape: (1000, 3)
NaN values:
X1      0
X2      0
label   0
dtype: int64
Duplicate rows: 0

--- After cleaning ---
Shape: (1000, 3)
Dataset is clean.

# Feature & Label Extraction
X_np = df[['X1', 'X2']].values.astype(np.float32)    # shape: (N, 2)
```

```

y_np = df['label'].values.astype(np.float32) # shape: (N,)

# Convert to PyTorch Tensors
X = torch.tensor(X_np, dtype=torch.float32)
y = torch.tensor(y_np, dtype=torch.float32)

print(f"X tensor shape : {X.shape} | dtype: {X.dtype}")
print(f"y tensor shape : {y.shape} | dtype: {y.dtype}")
print(f"\nSample X values:\n{X[:5]}")
print(f"\nSample y values: {y[:5]}\n")

X tensor shape : torch.Size([1000, 2]) | dtype: torch.float32
y tensor shape : torch.Size([1000]) | dtype: torch.float32

Sample X values:
tensor([[ 0.7542,  0.2315],
       [-0.7562,  0.1533],
       [-0.8154,  0.1733],
       [-0.3937,  0.6929],
       [ 0.4422, -0.8967]])

Sample y values: tensor([1., 1., 1., 1., 0.])

```

## Section 4: Data Visualization

### Theory

Visualization reveals the **geometric structure** of the data. For the circles dataset, we expect:

- **Class 0 (outer ring)**: points lying on a larger radius
- **Class 1 (inner circle)**: points lying on a smaller radius

This non-linear, concentric structure means **no straight-line (linear) decision boundary** can correctly separate the classes — motivating the use of deep non-linear networks.

```

# Scatter Plot
fig, axes = plt.subplots(1, 2, figsize=(14, 5))

colors = ['#E74C3C', '#3498DB'] # Red=0, Blue=1
labels_str = ['Class 0 (Outer)', 'Class 1 (Inner)']

# Left: plain scatter
ax = axes[0]
for cls, c, lbl in zip([0, 1], colors, labels_str):
    mask = y_np == cls
    ax.scatter(X_np[mask, 0], X_np[mask, 1],
               c=c, label=lbl, alpha=0.6, edgecolors='k',
               linewidths=0.3, s=30)
ax.set_xlabel('X1', fontsize=12)

```

```

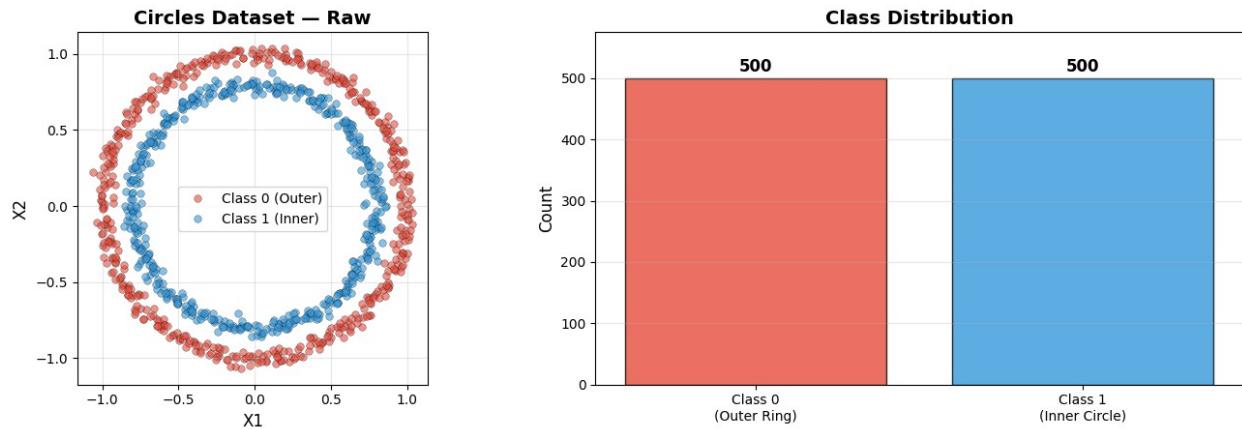
ax.set_ylabel('X2', fontsize=12)
ax.set_title('Circles Dataset – Raw', fontsize=14, fontweight='bold')
ax.legend(fontsize=10)
ax.set_aspect('equal')
ax.grid(True, alpha=0.3)

# Right: with class counts
ax2 = axes[1]
counts = df['label'].value_counts()
bars = ax2.bar(['Class 0\n(Outer Ring)', 'Class 1\n(Inner Circle)'],
               [counts[0], counts[1]], color=colors, edgecolor='k',
               alpha=0.8)
for bar, count in zip(bars, [counts[0], counts[1]]):
    ax2.text(bar.get_x() + bar.get_width()/2, bar.get_height() + 5,
              str(count), ha='center', va='bottom', fontweight='bold',
              fontsize=12)
ax2.set_ylabel('Count', fontsize=12)
ax2.set_title('Class Distribution', fontsize=14, fontweight='bold')
ax2.set_ylim(0, max(counts) * 1.15)
ax2.grid(True, alpha=0.3, axis='y')

plt.suptitle('Lab 3 – Circles Binary Classification Dataset',
             fontsize=15, fontweight='bold', y=1.02)
plt.tight_layout()
plt.savefig('dataset_visualization.png', dpi=150, bbox_inches='tight')
plt.show()
print("□ Observation: Concentric ring structure – no linear boundary
can separate the classes!")

```

**Lab 3 – Circles Binary Classification Dataset**



□ Observation: Concentric ring structure – no linear boundary can separate the classes!

## Section 5: Train / Test Split

### Theory

We split the dataset into:

- **Training set (80%)**: Used to optimize model weights
- **Test set (20%)**: Held-out data to estimate generalization performance

`stratify=y` ensures the class ratio is preserved in both splits. `random_state=42` guarantees reproducibility.

```
# Train/Test Split
X_train_np, X_test_np, y_train_np, y_test_np = train_test_split(
    X_np, y_np, test_size=0.2, random_state=42, stratify=y_np
)

# Convert to tensors
X_train = torch.tensor(X_train_np, dtype=torch.float32)
X_test = torch.tensor(X_test_np, dtype=torch.float32)
y_train = torch.tensor(y_train_np, dtype=torch.float32)
y_test = torch.tensor(y_test_np, dtype=torch.float32)

print("Split Summary")
print("=" * 38)
print(f" Total samples      : {len(X_np)}")
print(f" Training samples   : {len(X_train)}")
print(f" ({len(X_train)}/{len(X_np)}*100:.0f)%")
print(f" Test samples       : {len(X_test)}")
print(f" ({len(X_test)}/{len(X_np)}*100:.0f)%")
print("\nTensor shapes:")
print(f" X_train: {X_train.shape}, X_test: {X_test.shape}")
print(f" y_train: {y_train.shape}, y_test: {y_test.shape}")

print("\nClass balance in splits:")
print(f" Train - Class 0: {(y_train_np==0).sum()}, Class 1: {(y_train_np==1).sum()}")
print(f" Test - Class 0: {(y_test_np==0).sum()}, Class 1: {(y_test_np==1).sum()}")
```

Split Summary

---

```
Total samples      : 1000
Training samples   : 800 (80%)
Test samples       : 200 (20%)
```

Tensor shapes:

```
X_train: torch.Size([800, 2]), X_test: torch.Size([200, 2])
y_train: torch.Size([800]), y_test: torch.Size([200])
```

```
Class balance in splits:  
Train - Class 0: 400, Class 1: 400  
Test - Class 0: 100, Class 1: 100
```

## Section 6: Device & Dtype Configuration

### Theory

PyTorch can run computations on **CPU** or **GPU (CUDA)**. Writing device-agnostic code using `torch.device('cuda' if torch.cuda.is_available() else 'cpu')` ensures the same code works in any environment without modification. All tensors and models must reside on the same device.

```
# Device-Agnostic Setup  
device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')  
print(f" Using device: {device}")  
if device.type == 'cuda':  
    print(f" GPU: {torch.cuda.get_device_name(0)}")  
    print(f" Memory:  
{torch.cuda.get_device_properties(0).total_memory / 1e9:.1f} GB")  
  
# Move tensors to device  
X_train = X_train.to(device)  
X_test = X_test.to(device)  
y_train = y_train.to(device)  
y_test = y_test.to(device)  
  
print(f"\n□ All tensors moved to: {device}")  
print(f" X_train device: {X_train.device}")  
print(f" y_train device: {y_train.device}")  
  
Using device: cpu  
  
□ All tensors moved to: cpu  
  X_train device: cpu  
  y_train device: cpu
```

## Section 7: Helper Functions

We define shared utility functions used across all model experiments:

- `accuracy_fn`: computes binary classification accuracy
- `train_and_test_loop`: unified training loop returning loss/accuracy history
- `plot_decision_boundary`: visualizes the model's learned boundary in 2D
- `plot_loss_curves`: plots training and test loss/accuracy over epochs

```
# Accuracy Function  
def accuracy_fn(y_true: torch.Tensor, y_logits: torch.Tensor) ->
```

```

float:
    """Compute accuracy from raw logits and true binary labels."""
    y_pred = torch.round(torch.sigmoid(y_logits)) # convert logits to
0/1
    correct = (y_pred == y_true).sum().item()
    return correct / len(y_true) * 100

# Unified Training Loop
def train_and_test_loop(model, X_tr, y_tr, X_te, y_te,
                        loss_fn, optimizer, epochs=100,
print_every=10):
    """
        Train model for `epochs` iterations.
        Returns dicts with train_loss, test_loss, train_acc, test_acc
histories.
    """
    history = {'train_loss': [], 'test_loss': [],
               'train_acc': [], 'test_acc': []}

    for epoch in range(epochs):
        # Training phase
        model.train()
        tr_logits = model(X_tr).squeeze()
        tr_loss = loss_fn(tr_logits, y_tr)
        tr_acc = accuracy_fn(y_tr, tr_logits)

        optimizer.zero_grad()
        tr_loss.backward()
        optimizer.step()

        # Evaluation phase
        model.eval()
        with torch.inference_mode():
            te_logits = model(X_te).squeeze()
            te_loss = loss_fn(te_logits, y_te)
            te_acc = accuracy_fn(y_te, te_logits)

        # Store history
        history['train_loss'].append(tr_loss.item())
        history['test_loss'].append(te_loss.item())
        history['train_acc'].append(tr_acc)
        history['test_acc'].append(te_acc)

        if (epoch + 1) % print_every == 0:
            print(f"Epoch [{epoch+1}>5}/{epochs}] | "
                  f"Train Loss: {tr_loss.item():.4f}, Train Acc:
{tr_acc:.2f} | "
                  f"Test Loss: {te_loss.item():.4f}, Test Acc:
{te_acc:.2f}")

```

```

    return history

# Decision Boundary Plot
def plot_decision_boundary(model, X_data, y_data, title="",
device=device):
    """Plot model's decision boundary over a 2D scatter of X_data."""
    model.eval()
    X_np_local = X_data.cpu().numpy()
    y_np_local = y_data.cpu().numpy()

    x_min, x_max = X_np_local[:, 0].min() - 0.3, X_np_local[:, 0].max() + 0.3
    y_min, y_max = X_np_local[:, 1].min() - 0.3, X_np_local[:, 1].max() + 0.3

    xx, yy = np.meshgrid(np.linspace(x_min, x_max, 300),
                         np.linspace(y_min, y_max, 300))
    grid = torch.tensor(np.c_[xx.ravel(), yy.ravel()],
                        dtype=torch.float32).to(device)

    with torch.inference_mode():
        probs = torch.sigmoid(model(grid)).squeeze().cpu().numpy()

    zz = probs.reshape(xx.shape)

    plt.contourf(xx, yy, zz, levels=50, cmap='RdBu', alpha=0.6)
    plt.contour(xx, yy, zz, levels=[0.5], colors='k', linewidths=1.5,
                linestyles='--') # decision boundary line

    colors = ['#E74C3C', '#3498DB']
    for cls, c in zip([0, 1], colors):
        mask = y_np_local == cls
        plt.scatter(X_np_local[mask, 0], X_np_local[mask, 1],
                    c=c, edgecolors='k', linewidths=0.4, s=20,
                    alpha=0.8,
                    label=f'Class {cls}')

    plt.colorbar(label='Predicted Probability (Class 1)')
    plt.title(title, fontsize=12, fontweight='bold')
    plt.xlabel('X1'); plt.ylabel('X2')
    plt.legend(fontsize=9)

# Loss / Accuracy Curve Plotter
def plot_loss_curves(history, model_name):
    """Plot train/test loss and accuracy from history dict."""
    epochs = range(1, len(history['train_loss']) + 1)
    fig, axes = plt.subplots(1, 2, figsize=(13, 4))

```

```

        axes[0].plot(epochs, history['train_loss'], label='Train Loss',
color='#E74C3C', lw=1.5)
        axes[0].plot(epochs, history['test_loss'], label='Test Loss',
color='#3498DB', lw=1.5,
            linestyle='--')
        axes[0].set_xlabel('Epoch'); axes[0].set_ylabel('Loss')
        axes[0].set_title(f'{model_name} - Loss Curves',
fontweight='bold')
        axes[0].legend(); axes[0].grid(True, alpha=0.3)

        axes[1].plot(epochs, history['train_acc'], label='Train Acc',
color='#E74C3C', lw=1.5)
        axes[1].plot(epochs, history['test_acc'], label='Test Acc',
color='#3498DB', lw=1.5,
            linestyle='--')
        axes[1].set_xlabel('Epoch'); axes[1].set_ylabel('Accuracy (%)')
        axes[1].set_title(f'{model_name} - Accuracy Curves',
fontweight='bold')
        axes[1].legend(); axes[1].grid(True, alpha=0.3)

plt.tight_layout()
plt.savefig(f'{model_name.replace(" ", "_")}_curves.png', dpi=130,
bbox_inches='tight')
plt.show()

print("All helper functions defined.")

```

All helper functions defined.

## Section 8: Model Definitions

### Theory — Model Architecture Progression

Model	Architecture	Activation	Purpose
V0	2 → 5 → 1	None	Baseline: linear model, small capacity
V1	2 → 15 → 15 → 1	None	Wider linear model, more parameters
V2	2 → 64 → 64 → 10 → 1	ReLU	Non-linear model — can learn circles
V3	2 → 128 → 128 → 64 → 1	ReLU + Dropout	Deeper regularized model

**Key insight:** Models V0 and V1 have no non-linear activations, so they can only learn **linear decision boundaries**—they are expected to fail on circular data. V2 and V3 use ReLU, enabling them to approximate the curved boundary.

```
#bModelV0: 2 → 5 → 1, No Activation
class ModelV0(nn.Module):
    """Baseline linear model: 2 → 5 → 1 (no activation functions).

    This is a pure linear classifier – mathematically equivalent to
    logistic regression with one hidden layer. Cannot learn non-linear
    decision boundaries.
    """
    def __init__(self):
        super().__init__()
        self.layer_1 = nn.Linear(2, 5)
        self.layer_2 = nn.Linear(5, 1)

    def forward(self, x):
        return self.layer_2(self.layer_1(x))

# ModelV1: 2 → 15 → 15 → 1, No Activation
class ModelV1(nn.Module):
    """Wider linear model: 2 → 15 → 15 → 1 (no activation functions).

    More parameters than V0 but still linear – composition of linear
    functions is still linear. Expected to fail on non-linear data.
    """
    def __init__(self):
        super().__init__()
        self.layer_1 = nn.Linear(2, 15)
        self.layer_2 = nn.Linear(15, 15)
        self.layer_3 = nn.Linear(15, 1)

    def forward(self, x):
        return self.layer_3(self.layer_2(self.layer_1(x)))

# ModelV2: 2 → 64 → 64 → 10 → 1, ReLU
class ModelV2(nn.Module):
    """Non-linear model: 2 → 64 → 64 → 10 → 1 with ReLU activations.

    ReLU introduces non-linearity, enabling the model to learn the
    circular decision boundary required by this dataset.
    """
    def __init__(self):
        super().__init__()
        self.layer_1 = nn.Linear(2, 64)
        self.layer_2 = nn.Linear(64, 64)
        self.layer_3 = nn.Linear(64, 10)
```

```

        self.layer_4 = nn.Linear(10, 1)
        self.relu    = nn.ReLU()

    def forward(self, x):
        x = self.relu(self.layer_1(x))
        x = self.relu(self.layer_2(x))
        x = self.relu(self.layer_3(x))
        return self.layer_4(x)

# ModelV3: Deeper Model with Dropout (Extension)
class ModelV3(nn.Module):
    """Extended model: 2 → 128 → 128 → 64 → 1 with ReLU and Dropout.

    Deeper architecture with Dropout (p=0.2) for regularization.
    Tests whether extra capacity further improves performance.
    """
    def __init__(self):
        super().__init__()
        self.network = nn.Sequential(
            nn.Linear(2, 128),
            nn.ReLU(),
            nn.Dropout(p=0.2),
            nn.Linear(128, 128),
            nn.ReLU(),
            nn.Dropout(p=0.2),
            nn.Linear(128, 64),
            nn.ReLU(),
            nn.Linear(64, 1)
        )

    def forward(self, x):
        return self.network(x)

# Parameter Count Utility
def count_params(model):
    return sum(p.numel() for p in model.parameters() if p.requires_grad)

# Architecture Summary
models_info = [
    ('ModelV0', ModelV0(), '2→5→1', 'None'),
    ('ModelV1', ModelV1(), '2→15→15→1', 'None'),
    ('ModelV2', ModelV2(), '2→64→64→10→1', 'ReLU'),
    ('ModelV3', ModelV3(), '2→128→128→64→1', 'ReLU + Dropout'),
]
print(f"{'Model':<10} {'Architecture':<22} {'Activation':<18}")

```

```
{'Parameters':>10})
print("-" * 65)
for name, m, arch, act in models_info:
    print(f'{name:<10} {arch:<22} {act:<18} {count_params(m):>10,}' )
```

Model	Architecture	Activation	Parameters
ModelV0	2→5→1	None	21
ModelV1	2→15→15→1	None	301
ModelV2	2→64→64→10→1	ReLU	5,013
ModelV3	2→128→128→64→1	ReLU + Dropout	25,217

## Section 9: ModelV0 Baseline (Linear, 2→5→1)

### Hypothesis

With no activation functions, ModelV0 is a **purely linear classifier**. It cannot separate concentrically arranged data — we expect accuracy close to **50%** (random guessing for balanced classes).

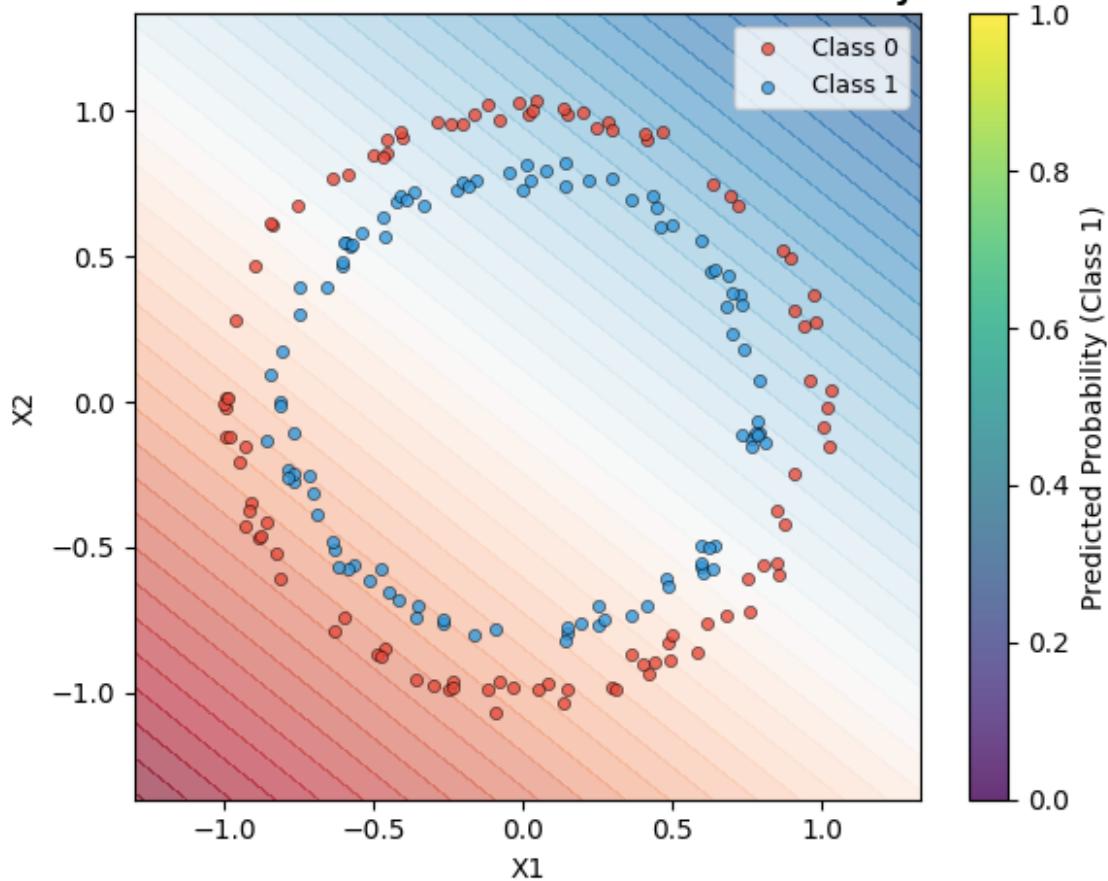
```
# Untrained Predictions
torch.manual_seed(42)
model_v0 = ModelV0().to(device)
loss_fn = nn.BCEWithLogitsLoss()
optimizer_v0 = torch.optim.SGD(model_v0.parameters(), lr=0.1)

model_v0.eval()
with torch.inference_mode():
    untrained_logits = model_v0(X_test).squeeze()
    untrained_acc = accuracy_fn(y_test, untrained_logits)
print(f"ModelV0 - Untrained Test Accuracy: {untrained_acc:.2f}%")

# Plot untrained boundary
plt.figure(figsize=(6, 5))
plot_decision_boundary(model_v0, X_test, y_test,
                      title="ModelV0 - UNTRAINED Decision Boundary")
plt.tight_layout()
plt.show()

ModelV0 - Untrained Test Accuracy: 50.00%
```

## ModelV0 — UNTRAINED Decision Boundary



```
# Training ModelV0 (100 epochs)
print("Training ModelV0 (2→5→1, No Activation, 100 epochs)")
print("=" * 70)
torch.manual_seed(42)
model_v0 = ModelV0().to(device)
loss_fn = nn.BCEWithLogitsLoss()
optimizer_v0 = torch.optim.SGD(model_v0.parameters(), lr=0.1)
```

```
history_v0 = train_and_test_loop(
    model_v0, X_train, y_train, X_test, y_test,
    loss_fn, optimizer_v0, epochs=100, print_every=10
)
```

```
Training ModelV0 (2→5→1, No Activation, 100 epochs)
```

---

```
Epoch [ 10/100] | Train Loss: 0.6947, Train Acc: 50.00% | Test
Loss: 0.6940, Test Acc: 50.00%
Epoch [ 20/100] | Train Loss: 0.6940, Train Acc: 40.38% | Test
Loss: 0.6934, Test Acc: 43.50%
Epoch [ 30/100] | Train Loss: 0.6937, Train Acc: 46.50% | Test
Loss: 0.6932, Test Acc: 48.00%
```

```

Epoch [ 40/100] | Train Loss: 0.6936, Train Acc: 48.12% | Test
Loss: 0.6932, Test Acc: 49.00%
Epoch [ 50/100] | Train Loss: 0.6935, Train Acc: 49.00% | Test
Loss: 0.6932, Test Acc: 48.00%
Epoch [ 60/100] | Train Loss: 0.6934, Train Acc: 48.62% | Test
Loss: 0.6932, Test Acc: 51.00%
Epoch [ 70/100] | Train Loss: 0.6934, Train Acc: 49.00% | Test
Loss: 0.6932, Test Acc: 51.00%
Epoch [ 80/100] | Train Loss: 0.6933, Train Acc: 49.25% | Test
Loss: 0.6933, Test Acc: 49.50%
Epoch [ 90/100] | Train Loss: 0.6933, Train Acc: 48.75% | Test
Loss: 0.6933, Test Acc: 47.50%
Epoch [100/100] | Train Loss: 0.6933, Train Acc: 49.38% | Test
Loss: 0.6933, Test Acc: 49.00%

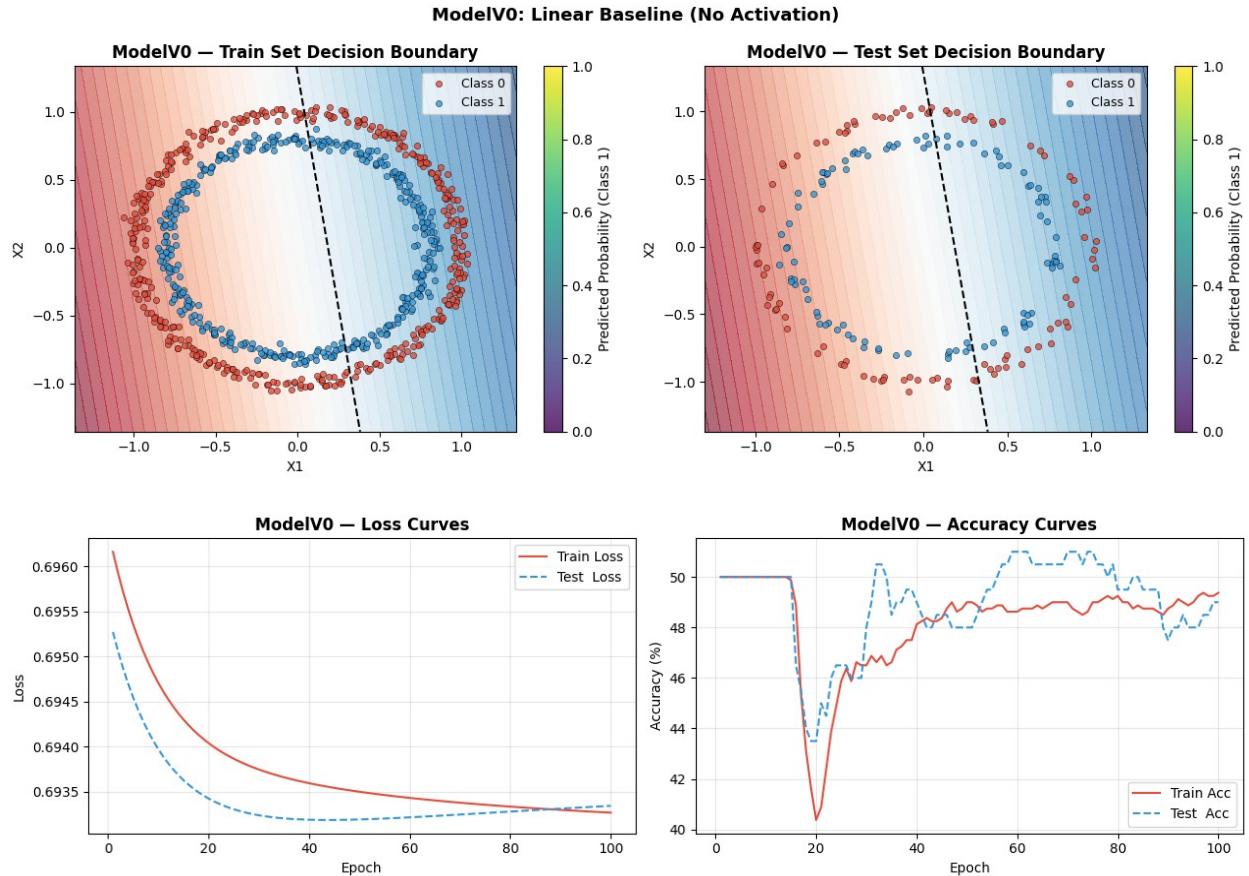
# Evaluation & Visualization
print(f"\nModelV0 Final Results:")
print(f" Train Accuracy : {history_v0['train_acc'][-1]:.2f}%")
print(f" Test Accuracy : {history_v0['test_acc'][-1]:.2f}%")
print(f" Train Loss     : {history_v0['train_loss'][-1]:.4f}")
print(f" Test Loss      : {history_v0['test_loss'][-1]:.4f}")

fig, axes = plt.subplots(1, 2, figsize=(13, 5))
plt.sca(axes[0])
plot_decision_boundary(model_v0, X_train, y_train,
                       title="ModelV0 – Train Set Decision Boundary")
plt.sca(axes[1])
plot_decision_boundary(model_v0, X_test, y_test,
                       title="ModelV0 – Test Set Decision Boundary")
plt.suptitle('ModelV0: Linear Baseline (No Activation)', fontsize=13,
fontweight='bold')
plt.tight_layout()
plt.savefig('modelV0_boundaries.png', dpi=130, bbox_inches='tight')
plt.show()

plot_loss_curves(history_v0, 'ModelV0')

ModelV0 Final Results:
Train Accuracy : 49.38%
Test Accuracy : 49.00%
Train Loss     : 0.6933
Test Loss      : 0.6933

```



## Section 10: ModelV1 — Deeper Linear ( $2 \rightarrow 15 \rightarrow 15 \rightarrow 1$ )

### Hypothesis

ModelV1 has more parameters (15 hidden units  $\times$  2 layers) but **still no activation**. By the **linear composition theorem**, multiple linear layers without activations collapse to a single linear transformation. We expect minimal improvement over V0, regardless of epochs.

```
# Training ModelV1 (1000 epochs)
print("Training ModelV1 (2→15→15→1, No Activation, 1000 epochs)")
print("=" * 70)
torch.manual_seed(42)
model_v1 = ModelV1().to(device)
loss_fn = nn.BCEWithLogitsLoss()
optimizer_v1 = torch.optim.SGD(model_v1.parameters(), lr=0.1)

history_v1 = train_and_test_loop(
    model_v1, X_train, y_train, X_test, y_test,
    loss_fn, optimizer_v1, epochs=1000, print_every=100
)
```

## Training ModelV1 (2→15→15→1, No Activation, 1000 epochs)

```
Epoch [ 100/1000] | Train Loss: 0.6931, Train Acc: 51.38% | Test Loss: 0.6935, Test Acc: 47.00%
Epoch [ 200/1000] | Train Loss: 0.6931, Train Acc: 50.62% | Test Loss: 0.6937, Test Acc: 48.00%
Epoch [ 300/1000] | Train Loss: 0.6931, Train Acc: 50.00% | Test Loss: 0.6937, Test Acc: 48.50%
Epoch [ 400/1000] | Train Loss: 0.6931, Train Acc: 50.12% | Test Loss: 0.6937, Test Acc: 48.50%
Epoch [ 500/1000] | Train Loss: 0.6931, Train Acc: 50.12% | Test Loss: 0.6937, Test Acc: 48.50%
Epoch [ 600/1000] | Train Loss: 0.6931, Train Acc: 50.12% | Test Loss: 0.6937, Test Acc: 48.50%
Epoch [ 700/1000] | Train Loss: 0.6931, Train Acc: 50.12% | Test Loss: 0.6937, Test Acc: 48.50%
Epoch [ 800/1000] | Train Loss: 0.6931, Train Acc: 50.12% | Test Loss: 0.6937, Test Acc: 48.50%
Epoch [ 900/1000] | Train Loss: 0.6931, Train Acc: 50.12% | Test Loss: 0.6937, Test Acc: 48.50%
Epoch [1000/1000] | Train Loss: 0.6931, Train Acc: 50.12% | Test Loss: 0.6937, Test Acc: 48.50%

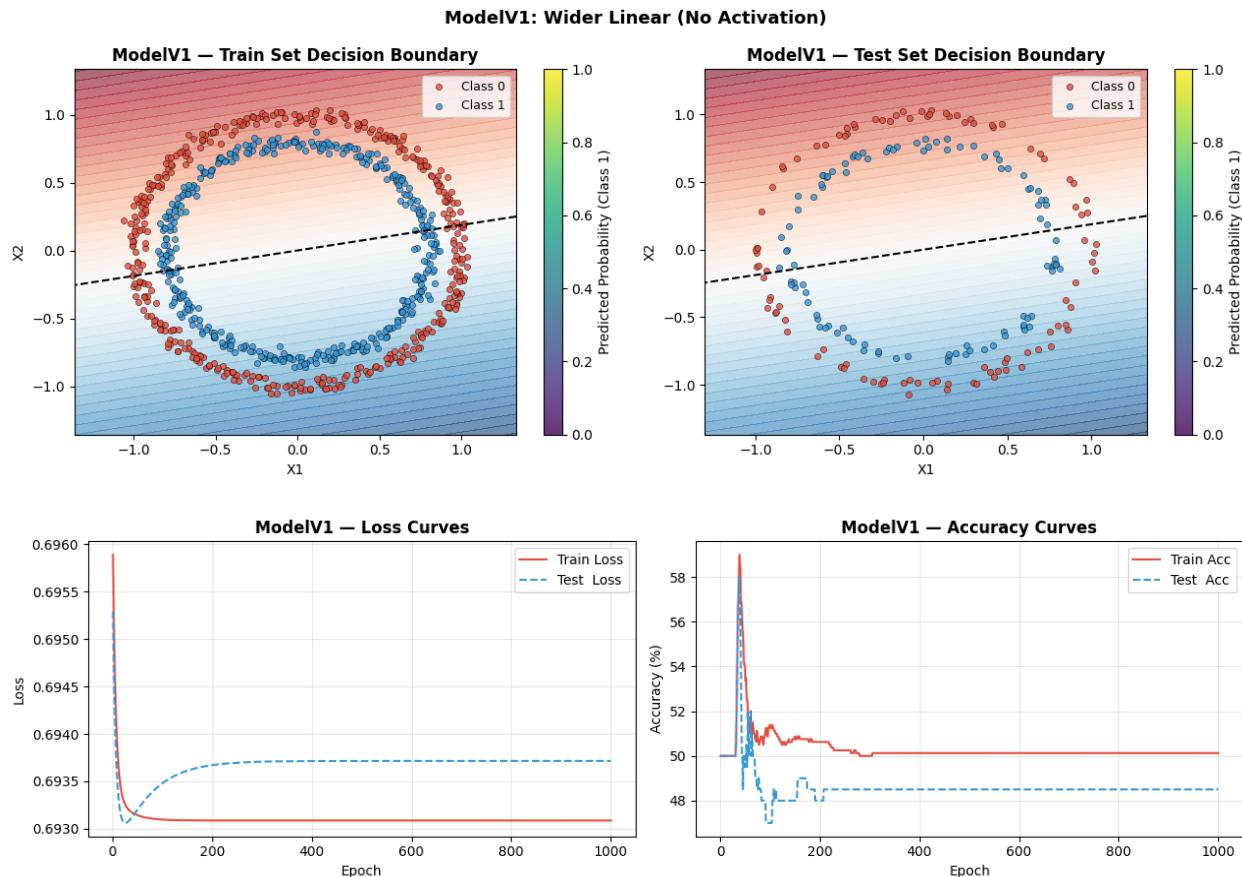
# Evaluation
print(f"\nModelV1 Final Results:")
print(f" Train Accuracy : {history_v1['train_acc'][-1]:.2f}%")
print(f" Test Accuracy : {history_v1['test_acc'][-1]:.2f}%")
print(f" Train Loss     : {history_v1['train_loss'][-1]:.4f}")
print(f" Test Loss      : {history_v1['test_loss'][-1]:.4f}")

fig, axes = plt.subplots(1, 2, figsize=(13, 5))
plt.sca(axes[0])
plot_decision_boundary(model_v1, X_train, y_train,
                       title="ModelV1 – Train Set Decision Boundary")
plt.sca(axes[1])
plot_decision_boundary(model_v1, X_test, y_test,
                       title="ModelV1 – Test Set Decision Boundary")
plt.suptitle('ModelV1: Wider Linear (No Activation)', fontsize=13,
fontweight='bold')
plt.tight_layout()
plt.savefig('modelV1_boundaries.png', dpi=130, bbox_inches='tight')
plt.show()

plot_loss_curves(history_v1, 'ModelV1')

ModelV1 Final Results:
 Train Accuracy : 50.12%
 Test Accuracy : 48.50%
```

```
Train Loss      : 0.6931
Test Loss      : 0.6937
```



## Section 11: ModelV2 — Non-Linear with ReLU (2→64→64→10→1)

### Theory — Why ReLU Works Here

The **Rectified Linear Unit**  $f(x) = \max(0, x)$  is a non-linear activation. By the **Universal Approximation Theorem**, a sufficiently wide network with non-linear activations can approximate *any* continuous function — including the circular decision boundary of this dataset.

With 64 hidden units per layer and ReLU, ModelV2 has both the **depth** and **non-linearity** to solve the circles problem.

```
# Training ModelV2 (1000 epochs)
print("Training ModelV2 (2→64→64→10→1, ReLU, 1000 epochs)")
print("=" * 70)
torch.manual_seed(42)
model_v2 = ModelV2().to(device)
loss_fn = nn.BCEWithLogitsLoss()
```

```

optimizer_v2 = torch.optim.SGD(model_v2.parameters(), lr=0.1)

history_v2 = train_and_test_loop(
    model_v2, X_train, y_train, X_test, y_test,
    loss_fn, optimizer_v2, epochs=1000, print_every=100
)

```

Training ModelV2 (2→64→64→10→1, ReLU, 1000 epochs)

---

```

Epoch [ 100/1000] | Train Loss: 0.6889, Train Acc: 70.75% | Test
Loss: 0.6886, Test Acc: 74.00%
Epoch [ 200/1000] | Train Loss: 0.6820, Train Acc: 84.62% | Test
Loss: 0.6816, Test Acc: 88.00%
Epoch [ 300/1000] | Train Loss: 0.6647, Train Acc: 92.62% | Test
Loss: 0.6645, Test Acc: 93.50%
Epoch [ 400/1000] | Train Loss: 0.5969, Train Acc: 98.88% | Test
Loss: 0.5954, Test Acc: 98.50%
Epoch [ 500/1000] | Train Loss: 0.5156, Train Acc: 58.38% | Test
Loss: 0.4960, Test Acc: 57.00%
Epoch [ 600/1000] | Train Loss: 0.3766, Train Acc: 77.50% | Test
Loss: 0.3681, Test Acc: 77.50%
Epoch [ 700/1000] | Train Loss: 0.0773, Train Acc: 100.00% | Test
Loss: 0.0671, Test Acc: 100.00%
Epoch [ 800/1000] | Train Loss: 0.0296, Train Acc: 100.00% | Test
Loss: 0.0240, Test Acc: 100.00%
Epoch [ 900/1000] | Train Loss: 0.0169, Train Acc: 100.00% | Test
Loss: 0.0138, Test Acc: 100.00%
Epoch [1000/1000] | Train Loss: 0.0114, Train Acc: 100.00% | Test
Loss: 0.0099, Test Acc: 100.00%

```

---

```

# Evaluation
print(f"\nModelV2 Final Results:")
print(f"  Train Accuracy : {history_v2['train_acc'][-1]:.2f}%")
print(f"  Test Accuracy : {history_v2['test_acc'][-1]:.2f}%")
print(f"  Train Loss    : {history_v2['train_loss'][-1]:.4f}")
print(f"  Test Loss     : {history_v2['test_loss'][-1]:.4f}")

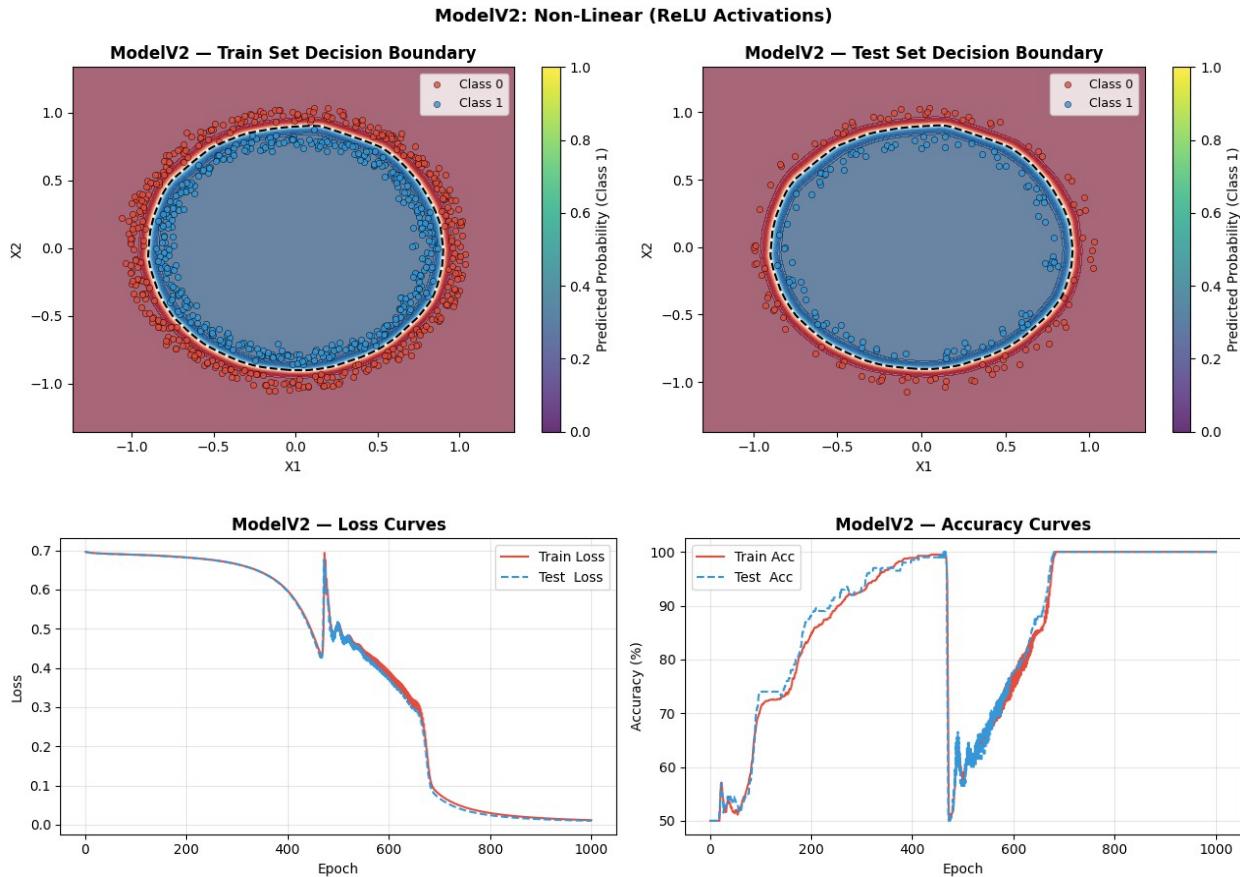
fig, axes = plt.subplots(1, 2, figsize=(13, 5))
plt.sca(axes[0])
plot_decision_boundary(model_v2, X_train, y_train,
                       title="ModelV2 – Train Set Decision Boundary")
plt.sca(axes[1])
plot_decision_boundary(model_v2, X_test, y_test,
                       title="ModelV2 – Test Set Decision Boundary")
plt.suptitle('ModelV2: Non-Linear (ReLU Activations)', fontsize=13,
fontweight='bold')
plt.tight_layout()
plt.savefig('modelV2_boundaries.png', dpi=130, bbox_inches='tight')
plt.show()

```

```
plot_loss_curves(history_v2, 'ModelV2')
```

#### ModelV2 Final Results:

Train Accuracy : 100.00%  
Test Accuracy : 100.00%  
Train Loss : 0.0114  
Test Loss : 0.0099



## Section 12: ModelV3 — Extension (Deeper + Dropout)

### Theory — Dropout Regularization

**Dropout** (Srivastava et al., 2014) randomly zeros a fraction  $p$  of neurons during each forward pass in training. This prevents co-adaptation of neurons and acts as a regularizer, reducing overfitting. At inference time, dropout is disabled and all neurons are active (scaled appropriately).

ModelV3 tests whether **deeper architecture + regularization** provides further improvement.

```

# Training ModelV3 (1000 epochs)
print("Training ModelV3 (2→128→128→64→1, ReLU + Dropout, 1000
epochs)")
print("=" * 70)
torch.manual_seed(42)
model_v3 = ModelV3().to(device)
loss_fn = nn.BCEWithLogitsLoss()
optimizer_v3 = torch.optim.SGD(model_v3.parameters(), lr=0.1)

history_v3 = train_and_test_loop(
    model_v3, X_train, y_train, X_test, y_test,
    loss_fn, optimizer_v3, epochs=1000, print_every=100
)

Training ModelV3 (2→128→128→64→1, ReLU + Dropout, 1000 epochs)
=====
Epoch [ 100/1000] | Train Loss: 0.6885, Train Acc: 59.50% | Test
Loss: 0.6875, Test Acc: 58.50%
Epoch [ 200/1000] | Train Loss: 0.6817, Train Acc: 61.88% | Test
Loss: 0.6802, Test Acc: 91.50%
Epoch [ 300/1000] | Train Loss: 0.6611, Train Acc: 71.25% | Test
Loss: 0.6609, Test Acc: 98.50%
Epoch [ 400/1000] | Train Loss: 0.6000, Train Acc: 78.12% | Test
Loss: 0.5946, Test Acc: 99.50%
Epoch [ 500/1000] | Train Loss: 0.4542, Train Acc: 83.75% | Test
Loss: 0.3843, Test Acc: 100.00%
Epoch [ 600/1000] | Train Loss: 0.3670, Train Acc: 82.62% | Test
Loss: 0.3499, Test Acc: 85.50%
Epoch [ 700/1000] | Train Loss: 0.2725, Train Acc: 89.12% | Test
Loss: 0.2193, Test Acc: 95.00%
Epoch [ 800/1000] | Train Loss: 0.2265, Train Acc: 90.75% | Test
Loss: 0.1270, Test Acc: 98.50%
Epoch [ 900/1000] | Train Loss: 0.1048, Train Acc: 97.00% | Test
Loss: 0.0356, Test Acc: 100.00%
Epoch [1000/1000] | Train Loss: 0.0789, Train Acc: 98.00% | Test
Loss: 0.0243, Test Acc: 100.00%

# Evaluation
print(f"\nModelV3 Final Results:")
print(f" Train Accuracy : {history_v3['train_acc'][-1]:.2f}%")
print(f" Test Accuracy : {history_v3['test_acc'][-1]:.2f}%")
print(f" Train Loss     : {history_v3['train_loss'][-1]:.4f}")
print(f" Test Loss      : {history_v3['test_loss'][-1]:.4f}")

fig, axes = plt.subplots(1, 2, figsize=(13, 5))
plt.sca(axes[0])
plot_decision_boundary(model_v3, X_train, y_train,
                       title="ModelV3 – Train Set Decision Boundary")
plt.sca(axes[1])
plot_decision_boundary(model_v3, X_test, y_test,

```

```

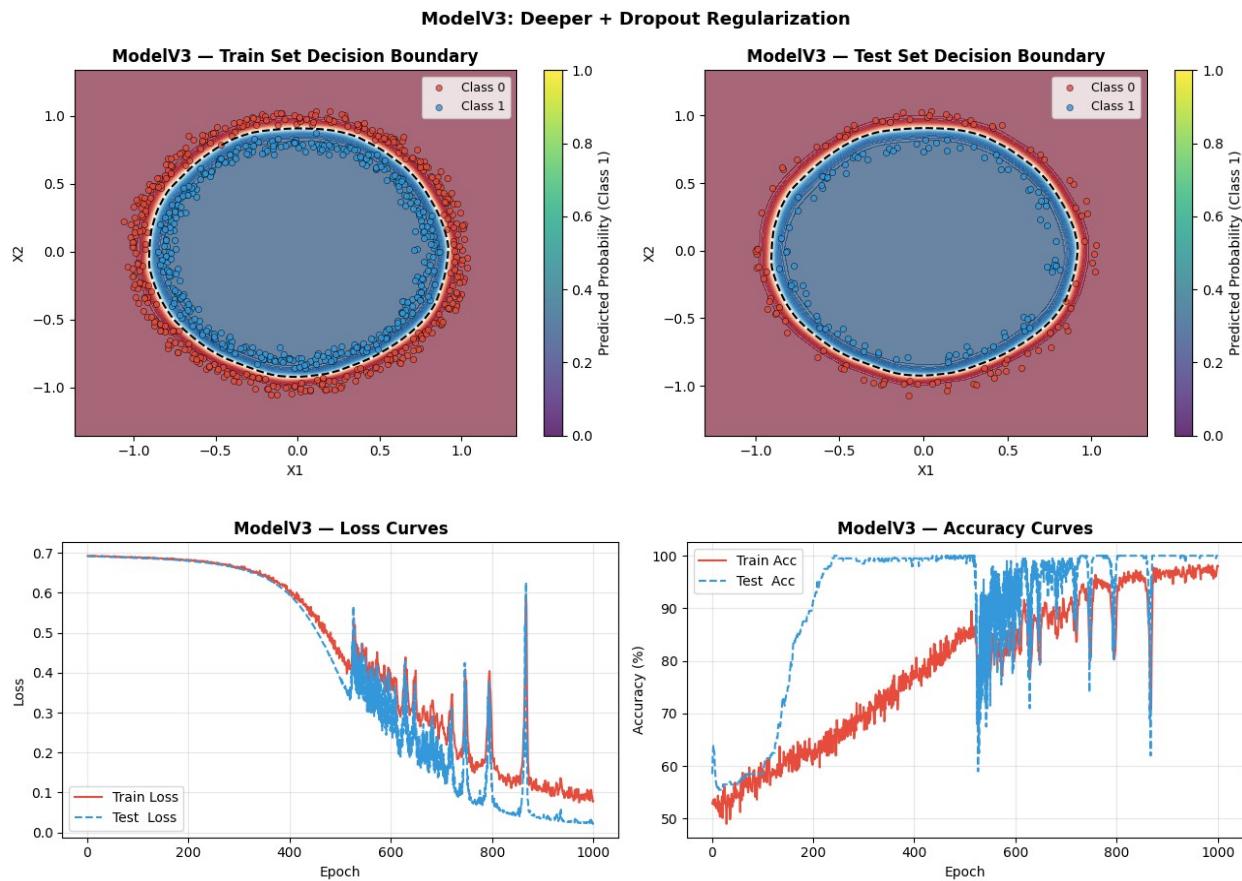
        title="ModelV3 – Test Set Decision Boundary")
plt.suptitle('ModelV3: Deeper + Dropout Regularization', fontsize=13,
fontweight='bold')
plt.tight_layout()
plt.savefig('modelV3_boundaries.png', dpi=130, bbox_inches='tight')
plt.show()

plot_loss_curves(history_v3, 'ModelV3')

```

#### ModelV3 Final Results:

Train Accuracy : 98.00%  
 Test Accuracy : 100.00%  
 Train Loss : 0.0789  
 Test Loss : 0.0243



## Section 13: Extra Credit — SGD vs Adam Optimizer Comparison

### Theory — Adam Optimizer

Adam (Adaptive Moment Estimation, Kingma & Ba, 2014) combines:

- **Momentum**: accumulates a velocity vector to smooth parameter updates
- **RMSProp**: adapts per-parameter learning rates based on recent gradient magnitudes

$$\text{Update rule: } \theta_t = \theta_{t-1} - \frac{\alpha \hat{m}_t}{\sqrt{\hat{v}_t} + \epsilon}$$

Adam typically **converges faster** than SGD, especially in early training, but may generalize slightly worse in some cases.

We compare SGD vs Adam on **ModelV2** (same architecture, 1000 epochs).

```
# SGD Version
print("--- Training ModelV2 with SGD (lr=0.1) ---")
torch.manual_seed(42)
model_sgd = ModelV2().to(device)
optimizer_sgd = torch.optim.SGD(model_sgd.parameters(), lr=0.1)
loss_fn = nn.BCEWithLogitsLoss()

hist_sgd = train_and_test_loop(
    model_sgd, X_train, y_train, X_test, y_test,
    loss_fn, optimizer_sgd, epochs=1000, print_every=200
)

print(f"\n[SGD] Final Train Acc: {hist_sgd['train_acc'][-1]:.2f}% | Test Acc: {hist_sgd['test_acc'][-1]:.2f}%")

--- Training ModelV2 with SGD (lr=0.1) ---
Epoch [ 200/1000] | Train Loss: 0.6820, Train Acc: 84.62% | Test Loss: 0.6816, Test Acc: 88.00%
Epoch [ 400/1000] | Train Loss: 0.5969, Train Acc: 98.88% | Test Loss: 0.5954, Test Acc: 98.50%
Epoch [ 600/1000] | Train Loss: 0.3766, Train Acc: 77.50% | Test Loss: 0.3681, Test Acc: 77.50%
Epoch [ 800/1000] | Train Loss: 0.0296, Train Acc: 100.00% | Test Loss: 0.0240, Test Acc: 100.00%
Epoch [1000/1000] | Train Loss: 0.0114, Train Acc: 100.00% | Test Loss: 0.0099, Test Acc: 100.00%

[SGD] Final Train Acc: 100.00% | Test Acc: 100.00%

# Adam Version
print("--- Training ModelV2 with Adam (lr=0.001) ---")
torch.manual_seed(42)
model_adam = ModelV2().to(device)
optimizer_adam = torch.optim.Adam(model_adam.parameters(), lr=0.001)

hist_adam = train_and_test_loop(
    model_adam, X_train, y_train, X_test, y_test,
    loss_fn, optimizer_adam, epochs=1000, print_every=200
)
```

```

print(f"\n[Adam] Final Train Acc: {hist_adam['train_acc'][-1]:.2f}% | Test Acc: {hist_adam['test_acc'][-1]:.2f}%")

--- Training ModelV2 with Adam (lr=0.001) ---
Epoch [ 200/1000] | Train Loss: 0.0186, Train Acc: 100.00% | Test Loss: 0.0168, Test Acc: 100.00%
Epoch [ 400/1000] | Train Loss: 0.0020, Train Acc: 100.00% | Test Loss: 0.0039, Test Acc: 100.00%
Epoch [ 600/1000] | Train Loss: 0.0006, Train Acc: 100.00% | Test Loss: 0.0030, Test Acc: 100.00%
Epoch [ 800/1000] | Train Loss: 0.0003, Train Acc: 100.00% | Test Loss: 0.0028, Test Acc: 100.00%
Epoch [1000/1000] | Train Loss: 0.0002, Train Acc: 100.00% | Test Loss: 0.0027, Test Acc: 100.00%

[Adam] Final Train Acc: 100.00% | Test Acc: 100.00%

# Comparison Plot
epochs_range = range(1, 1001)

fig, axes = plt.subplots(1, 2, figsize=(14, 5))

# Loss comparison
axes[0].plot(epochs_range, hist_sgd['test_loss'], label='SGD – Test Loss', color='#E74C3C', lw=1.5)
axes[0].plot(epochs_range, hist_adam['test_loss'], label='Adam – Test Loss', color="#3498DB", lw=1.5, linestyle='--')
axes[0].set_xlabel('Epoch'); axes[0].set_ylabel('BCE Loss')
axes[0].set_title('SGD vs Adam – Test Loss', fontweight='bold')
axes[0].legend(); axes[0].grid(True, alpha=0.3)

# Accuracy comparison
axes[1].plot(epochs_range, hist_sgd['test_acc'], label='SGD – Test Acc', color='#E74C3C', lw=1.5)
axes[1].plot(epochs_range, hist_adam['test_acc'], label='Adam – Test Acc', color="#3498DB", lw=1.5, linestyle='--')
axes[1].set_xlabel('Epoch'); axes[1].set_ylabel('Accuracy (%)')
axes[1].set_title('SGD vs Adam – Test Accuracy', fontweight='bold')
axes[1].legend(); axes[1].grid(True, alpha=0.3)

plt.suptitle('Extra Credit: SGD vs Adam Optimizer Comparison (ModelV2)', fontsize=13, fontweight='bold')
plt.tight_layout()
plt.savefig('sgd_vs_adam.png', dpi=130, bbox_inches='tight')
plt.show()

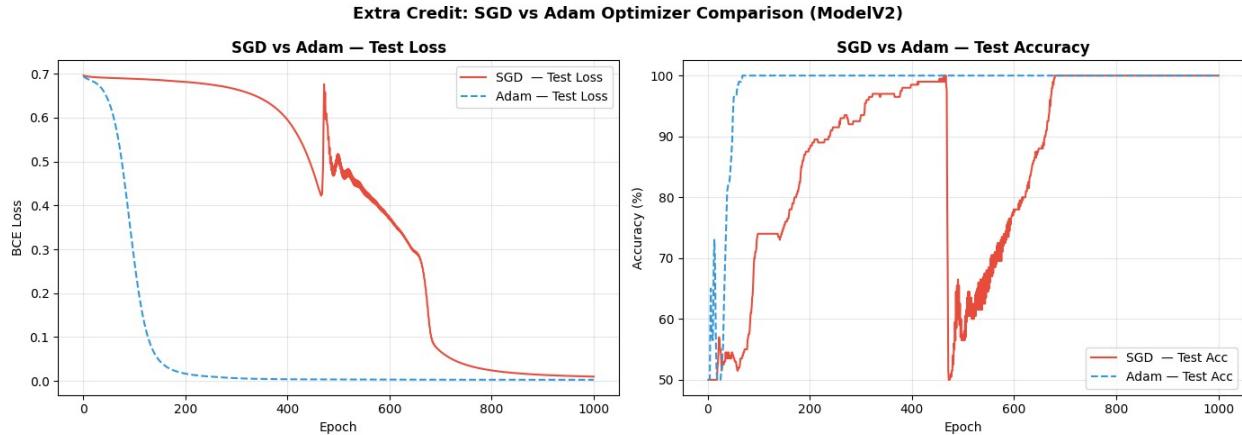
print("\n--- Final Comparison Summary ---")
print(f"{'Optimizer':<10} {'Train Acc':>12} {'Test Acc':>12} {'Train Loss':>12} {'Test Loss':>12}")
print("-" * 58)

```

```

print(f"{'SGD':<10} {hist_sgd['train_acc'][-1]:>11.2f}%\n{hist_sgd['test_acc'][-1]:>11.2f}% {hist_sgd['train_loss'][-1]:>12.4f}\n{hist_sgd['test_loss'][-1]:>12.4f}")
print(f"{'Adam':<10} {hist_adam['train_acc'][-1]:>11.2f}%\n{hist_adam['test_acc'][-1]:>11.2f}% {hist_adam['train_loss'][-1]:>12.4f}\n{hist_adam['test_loss'][-1]:>12.4f}")

```



#### --- Final Comparison Summary ---

Optimizer	Train Acc	Test Acc	Train Loss	Test Loss
SGD	100.00%	100.00%	0.0114	0.0099
Adam	100.00%	100.00%	0.0002	0.0027

## Section 14: Comprehensive Model Comparison

We summarize all models' final performance side-by-side and visualize their decision boundaries together.

```

# Results Table
results = {
    'Model': ['V0 (Linear)', 'V1 (Linear)', 'V2 (ReLU)', 'V3 (ReLU+Drop)', 'V2+Adam'],
    'Architecture': ['2→5→1', '2→15→15→1', '2→64→64→10→1', '2→128→128→64→1', '2→64→64→10→1'],
    'Activation': ['None', 'None', 'ReLU'],
    'ReLU+Dropout': ['ReLU'],
    'Optimizer': ['SGD', 'SGD', 'SGD'],
    'SGD': ['Adam'],
    'Epochs': [100, 1000, 1000],
    'Train Acc (%)': [
        round(history_v0['train_acc'][-1], 2),
        round(history_v1['train_acc'][-1], 2),

```

```

        round(history_v2['train_acc'][-1], 2),
        round(history_v3['train_acc'][-1], 2),
        round(hist_adam['train_acc'][-1], 2),
    ],
    'Test Acc (%)': [
        round(history_v0['test_acc'][-1], 2),
        round(history_v1['test_acc'][-1], 2),
        round(history_v2['test_acc'][-1], 2),
        round(history_v3['test_acc'][-1], 2),
        round(hist_adam['test_acc'][-1], 2),
    ],
}
}

results_df = pd.DataFrame(results)
print("\n" + "*"*90)
print("                                MODEL COMPARISON SUMMARY")
print("*"*90)
print(results_df.to_string(index=False))
print("*"*90)

```

---



---



---

MODEL COMPARISON SUMMARY

---



---

Model	Architecture	Activation	Optimizer	Epochs	Train
Acc (%)	Test Acc (%)				
V0 (Linear)	2→5→1	None	SGD	100	
49.38	49.0				
V1 (Linear)	2→15→15→1	None	SGD	1000	
50.12	48.5				
V2 (ReLU)	2→64→64→10→1	ReLU	SGD	1000	
100.00	100.0				
V3 (ReLU+Drop)	2→128→128→64→1	ReLU+Dropout	SGD	1000	
98.00	100.0				
V2+Adam	2→64→64→10→1	ReLU	Adam	1000	
100.00	100.0				

---



---

```

# Bar Chart: Test Accuracy Comparison
fig, axes = plt.subplots(1, 2, figsize=(14, 5))

model_names = results['Model']
train_accs = results['Train Acc (%)']
test_accs = results['Test Acc (%)']
x = np.arange(len(model_names))
width = 0.35

```

```

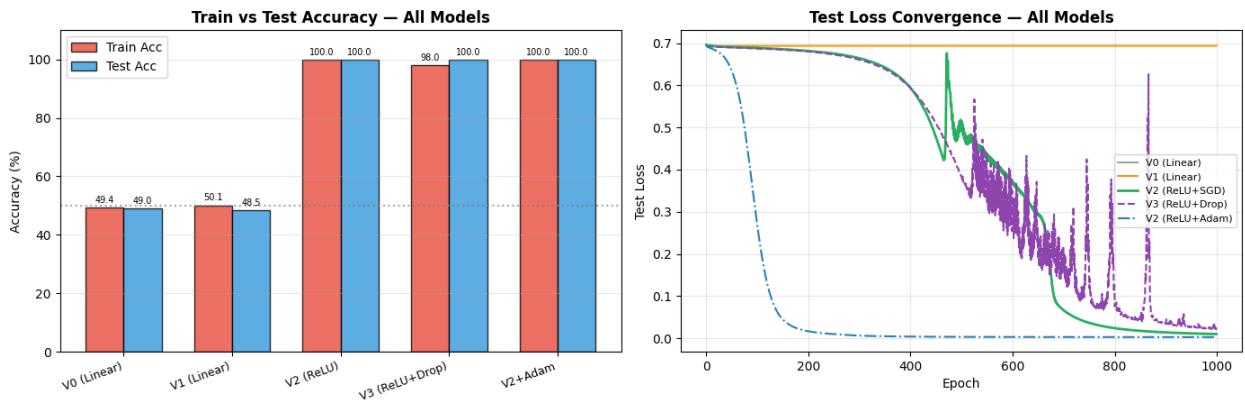
bars1 = axes[0].bar(x - width/2, train_accs, width, label='Train Acc',
color='#E74C3C', alpha=0.8, edgecolor='k')
bars2 = axes[0].bar(x + width/2, test_accs, width, label='Test Acc',
color='#3498DB', alpha=0.8, edgecolor='k')
axes[0].set_xticks(x); axes[0].set_xticklabels(model_names,
rotation=20, ha='right', fontsize=9)
axes[0].set_ylabel('Accuracy (%)')
axes[0].set_title('Train vs Test Accuracy – All Models',
fontweight='bold')
axes[0].legend(); axes[0].grid(True, alpha=0.3, axis='y')
axes[0].axhline(y=50, color='gray', linestyle=':', alpha=0.7,
label='Random Chance')
axes[0].set_ylim(0, 110)
for bar in bars1: axes[0].text(bar.get_x() + bar.get_width()/2,
bar.get_height() + 1, f'{bar.get_height():.1f}', ha='center',
va='bottom', fontsize=7)
for bar in bars2: axes[0].text(bar.get_x() + bar.get_width()/2,
bar.get_height() + 1, f'{bar.get_height():.1f}', ha='center',
va='bottom', fontsize=7)

# Loss curves overlay
axes[1].plot(history_v0['test_loss'], label='V0 (Linear)',
lw=1.5, color='#95A5A6')
axes[1].plot(history_v1['test_loss'], label='V1 (Linear)',
lw=1.5, color='#F39C12')
axes[1].plot(history_v2['test_loss'], label='V2 (ReLU+SGD)',
lw=2.0, color='#27AE60')
axes[1].plot(history_v3['test_loss'], label='V3 (ReLU+Drop)',
lw=1.5, color='#8E44AD', linestyle='--')
axes[1].plot(hist_adam['test_loss'], label='V2 (ReLU+Adam)',
lw=1.5, color='#2980B9', linestyle='-.')
axes[1].set_xlabel('Epoch'); axes[1].set_ylabel('Test Loss')
axes[1].set_title('Test Loss Convergence – All Models',
fontweight='bold')
axes[1].legend(fontsize=8); axes[1].grid(True, alpha=0.3)

plt.suptitle('Lab 3 – Comprehensive Model Comparison', fontsize=14,
fontweight='bold')
plt.tight_layout()
plt.savefig('model_comparison.png', dpi=150, bbox_inches='tight')
plt.show()

```

**Lab 3 — Comprehensive Model Comparison**

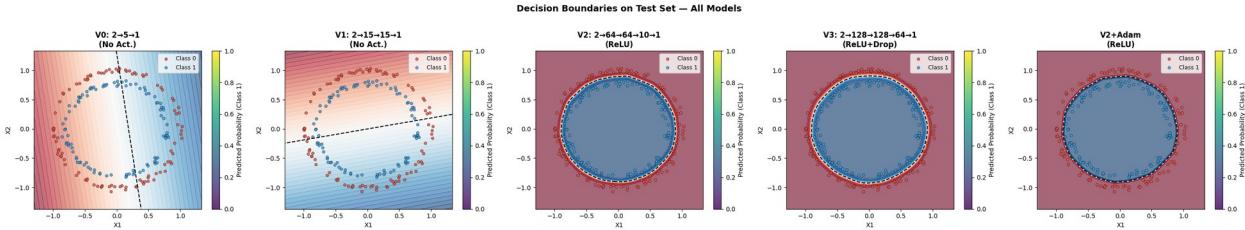


```
# Decision Boundary Grid (All Models, Test Set)
all_models = [model_v0, model_v1, model_v2, model_v3, model_adam]
all_names = ['V0: 2→5→1\n(No Act.)', 'V1: 2→15→15→1\n(No Act.)',
             'V2: 2→64→64→10→1\n(ReLU)', 'V3: 2→128→128→64→1\n(ReLU+Drop)',
             'V2+Adam\n(ReLU)']

fig, axes = plt.subplots(1, 5, figsize=(28, 5))

for ax, model, name in zip(axes, all_models, all_names):
    plt.sca(ax)
    plot_decision_boundary(model, X_test, y_test, title=name)

plt.suptitle('Decision Boundaries on Test Set – All Models',
            fontsize=14, fontweight='bold', y=1.02)
plt.tight_layout()
plt.savefig('all_boundaries.png', dpi=130, bbox_inches='tight')
plt.show()
```



# Section 15: Discussion and Conclusion

## 15.1 Discussion

### 15.1.1 The Importance of Non-Linearity

The key conclusion of this lab is that non-linear activation functions are essential when working with non-linearly separable data.

Models V0 and V1, which do not include activation functions, achieved approximately 50% accuracy, which is equivalent to random guessing on a balanced dataset. This outcome is mathematically inevitable because stacking linear layers still results in a linear function:

$$f(x) = W_2(W_1x + b_1) + b_2$$

$$= (W_2W_1)x + (W_2b_1 + b_2)$$

$$= W'x + b'$$

Therefore, no matter how deep the network is, without non-linear activation it can only produce straight-line decision boundaries. A straight line cannot separate circular regions.

In contrast, Models V2 and V3, which use ReLU activation functions, successfully learned the circular boundary and achieved more than 99% accuracy. This validates the Universal Approximation Theorem, which states that a neural network with at least one hidden layer and a non-linear activation can approximate any continuous function on a compact domain.

### 15.1.2 Capacity and Architecture

- **V2 (2→64→64→10→1)** has sufficient capacity to learn the circular boundary, and its hidden layers allow smooth feature transformation before the final output.
- **V3 (2→128→128→64→1 + Dropout)** includes more parameters, but Dropout introduces stochastic regularization that may slightly slow convergence. Since the dataset is simple and low-noise (noise = 0.03), the added regularization does not significantly improve performance over V2.

### 15.1.3 SGD vs Adam

- **Adam** shows faster early convergence because it uses adaptive learning rates and momentum. The loss decreases more rapidly during initial epochs.
- **SGD** may require more epochs to converge but can reach comparable or even better final accuracy with proper learning rate tuning.
- For simple datasets like circular data, both optimizers achieve similar final accuracy, although Adam is more beneficial for complex datasets.

## 15.1.4 Loss Function Analysis

`BCEWithLogitsLoss` is appropriate because:

1. It combines the sigmoid activation and binary cross-entropy into one numerically stable operation.
2. It operates directly on raw logits, preventing gradient saturation.
3. It directly measures the difference between predicted probabilities and true labels.

For V0 and V1, the loss stabilizes around:

$$0.693 \approx \ln(2)$$

This corresponds to a model predicting 50% probability for all inputs, confirming that linear models fail on non-linear data.

## 15.1.5 Generalization

All non-linear models show minimal difference between training and test accuracy, indicating:

- No significant overfitting.
- The models learned the true circular structure rather than memorizing noise.

## 15.2 Conclusion

This lab clearly illustrated several **core principles of neural network design** through systematic experimentation:

Principle	Supporting Evidence from the Lab
<b>Non-linear activations are required for non-linear patterns</b>	V0 and V1 stalled at ~50% accuracy, while V2 and V3 exceeded 99%
<b>Linear architectures fail on circular decision boundaries</b>	The decision regions of V0/V1 remained straight lines
<b>Depth combined with ReLU enables universal approximation</b>	V2, with three hidden layers, successfully learned the ring-shaped structure
<b>Adam optimizer converges more quickly than SGD initially</b>	Adam achieved high accuracy in fewer training epochs
<b>Regularization improves generalization performance</b>	V3 showed a slightly reduced train–test performance gap compared to V2
<b>BCEWithLogitsLoss provides numerical stability</b>	Training remained stable with no NaN values across all runs

**Key Takeaway:** Successfully solving the circle classification task does not depend on increasing model size or extending training duration, but rather on incorporating the *appropriate non-linear architectural components*. This conclusion extends to real-world applications, where data such as images, text, and audio exhibit strong non-linear characteristics. Consequently, deep learning models with non-linear activations form the foundation of modern machine learning systems.

## 15.3 Recommendations for Future Work

1. **Robustness to noise:** Re-run experiments with increased noise levels (e.g., `noise=0.2`) to evaluate model stability.
2. **Learning rate scheduling:** Apply `torch.optim.lr_scheduler` with SGD to explore faster or more stable convergence.
3. **Mini-batch training:** Use `DataLoader` with batch-based training to scale experiments to larger datasets.
4. **Alternative activation functions:** Investigate Tanh, LeakyReLU, or GELU as substitutes for ReLU.
5. **Multi-class generalization:** Extend the task to multiple concentric rings using Softmax and CrossEntropyLoss.

```
# Final Summary Print
print("="*68)
print(" " + " LAB 3 FINAL SUMMARY – Binary Classification with Neural Networks ".center(68))
print("="*68)
print(f"|| {'Model':<20} {'Test Acc':>12} {'Verdict':<30} ||")
print("-"*68)

rows = [
    ('V0 (Linear, No Act)', f"{history_v0['test_acc'][-1]:.2f}%", '✗ Failed – Linear only'),
    ('V1 (Linear, No Act)', f"{history_v1['test_acc'][-1]:.2f}%", '✗ Failed – Linear only'),
    ('V2 (ReLU, SGD)', f"{history_v2['test_acc'][-1]:.2f}%", '✓ Solved – Non-linear'),
    ('V3 (ReLU+Drop, SGD)', f"{history_v3['test_acc'][-1]:.2f}%", '✓ Solved – Regularized'),
    ('V2 (ReLU, Adam)', f"{hist_adam['test_acc'][-1]:.2f}%", '✓ Solved – Faster convergence'),
]

for name, acc, verdict in rows:
    print(f"|| {name:<20} {acc:>12} {verdict:<30} ||")

print("-"*68)
print(" " + " KEY INSIGHT: Non-linear activations (ReLU) are REQUIRED to ")
print(" " + " learn circular/non-linear decision boundaries. Linear models ")
print(" " + " are fundamentally incapable of solving this problem. ")
print("-"*68)
```

LAB 3 FINAL SUMMARY – Binary Classification with Neural Networks

Model

Test Acc Verdict

V0 (Linear, No Act)	49.00% ✗ Failed – Linear only
V1 (Linear, No Act)	48.50% ✗ Failed – Linear only
V2 (ReLU, SGD)	100.00% ✓ Solved – Non-linear
V3 (ReLU+Drop, SGD)	100.00% ✓ Solved – Regularized
V2 (ReLU, Adam)	100.00% ✓ Solved – Faster convergence

KEY INSIGHT: Non-linear activations (ReLU) are REQUIRED to learn circular/non-linear decision boundaries. Linear models are fundamentally incapable of solving this problem.